

School of Computer Science
UNIVERSITY OF PETROLEUM AND ENERGY STUDIES
DEHRADUN, UTTARAKHAND



Container Orchestration and Infrastructure Automation
ContactHub: Contacts Management System
Project Report
(2023-2024)
for
5th Semester

SUBMITTED

TO:

Ms.Avita Katal
Assistant Professor(SG)
School of Computer Science

SUBMITTED BY:

Sanaya Bhardwaj
SAP ID: 500091835
Roll no.: R2142210690

Part 2: Deployment and Scaling

PART 1 Recap:

ContactHub-The Contact Manager Microservices project leverages the power of Spring Boot, a robust Java-based framework, to create a scalable and highly maintainable solution for contact management.

Microservices are defined with clear purposes and communication methods, justified for scalability. Docker is chosen for containerization, ensuring portability and isolation.

Microservices seamlessly interact through APIs, emphasizing integration. Git is effectively used for version control.

Microservices created for the project:

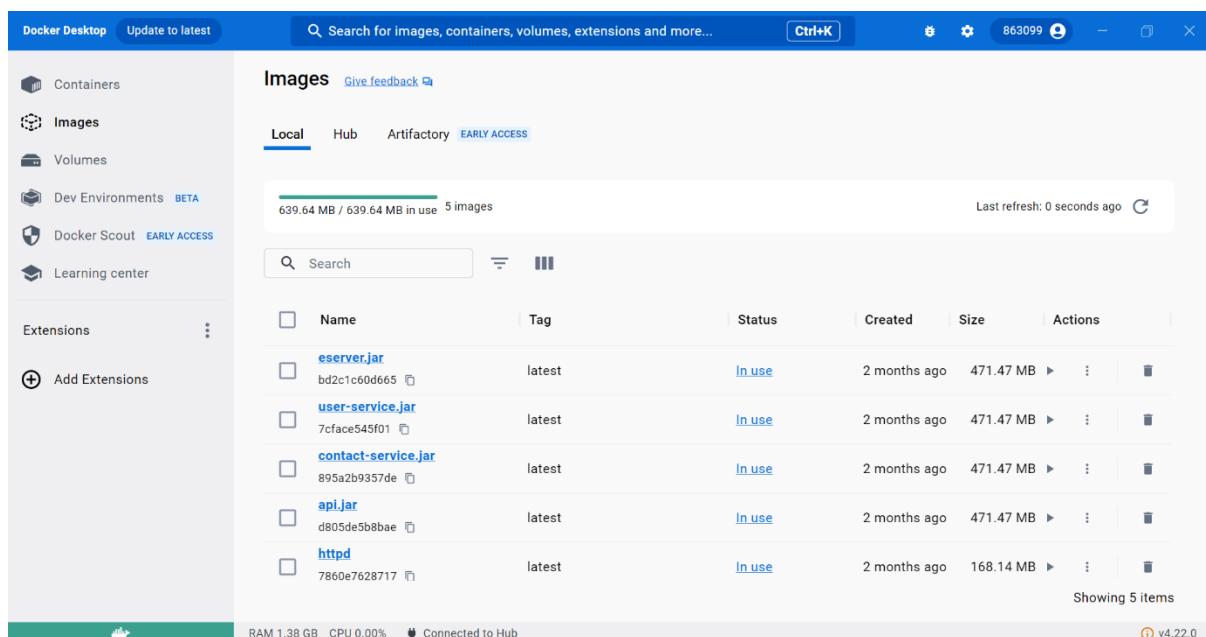
API Gateway: Serves as the entry point for client requests. It routes requests to the appropriate microservices based on the request path or URL.

Eureka (Service Registry): Microservices register themselves with Eureka upon startup. It maintains a registry of all available services and their network locations.

User Service: Manages user-related data. The API Gateway forwards user-related requests to this service.

Contact Service: Handles contact-related data. The API Gateway routes contact-related requests to this service.

Docker images created for Containerization:



Cloud Platform Selection

Justify the choice of a specific cloud platform (e.g., AWS, Azure, Google Cloud) based on its suitability for the project.

For ContactHub, the chosen cloud platform based on its suitability is AWS.

Amazon Web Services (AWS):

- Java Ecosystem Support:
 - AWS has excellent support for Java applications, making it well-suited for a project built on the Spring Boot framework.
 - AWS offers services like AWS Elastic Beanstalk and AWS Lambda, which are particularly useful for deploying and managing Java applications.
- Microservices Architecture:
 - AWS provides a range of services that support microservices architecture, such as Amazon ECS (Elastic Container Service) and AWS Lambda.
 - Integration with AWS API Gateway can help in creating and managing APIs for the microservices.
- Scalability and Flexibility:
 - AWS provides scalable solutions through services like Amazon EC2 for virtual servers and Amazon RDS for managed databases.
 - Auto Scaling features can automatically adjust capacity based on demand.
- Extensive Service Offering:
 - AWS has a comprehensive set of services, including databases (Amazon DynamoDB, Amazon RDS), storage (Amazon S3), and messaging (Amazon SQS), which are crucial for a microservices architecture.

Container Orchestration

Utilize a container orchestration platform (e.g., Kubernetes) to manage and deploy the microservices.

Installing Minikube:

What & Why Minikube?

minikube is local Kubernetes, focusing on making it easy to learn and develop for Kubernetes. All we need is Docker (or similarly compatible) container or a Virtual Machine environment, and Kubernetes

is a single command away: minikube start

```
Command Prompt - minikube
Microsoft Windows [Version 10.0.22621.2428]
(c) Microsoft Corporation. All rights reserved.

C:\Users\HP-PC>minikube start
W1117 00:47:52.297893 18068 main.go:291] Unable to resolve the current Docker CLI context "default": context "default"
: context not found: open C:\Users\HP-PC\.docker\contexts\meta\37a8eec1ce19687d132fe29051dca629d164e2c4958ba141d5f4133a3
3f0688f\meta.json: The system cannot find the path specified.
* minikube v1.32.0 on Microsoft Windows 11 Home Single Language 10.0.22621.2428 Build 22621.2428
* Automatically selected the docker driver. Other choices: virtualbox, ssh
* Using Docker Desktop driver with root privileges
* Starting control plane node minikube in cluster minikube
* Pulling base image ...
* Downloading Kubernetes v1.28.3 preload ...
  > gcr.io/k8s-minikube/kicbase...: 339.53 MiB / 453.90 MiB 74.80% 9.35 MiB/s
```

Deploying Microservices with Kubernetes:

- “docker-k8s.yaml” for contact service deployment

```
ApigatewayAp... ContactServ... Dockerfile application.yml docker-k8s.yaml Dockerfile
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: contactservice
5 spec:
6   replicas: 3
7   selector:
8     matchLabels:
9       app: contactservice
10  template:
11    metadata:
12      labels:
13        app: contactservice
14    spec:
15      containers:
16      - name: contactservice
17        image: 863099/springboot-image:latest
18        ports:
19        - containerPort: 3000
20
21
```

- “docker-k8s.yaml” for user service deployment

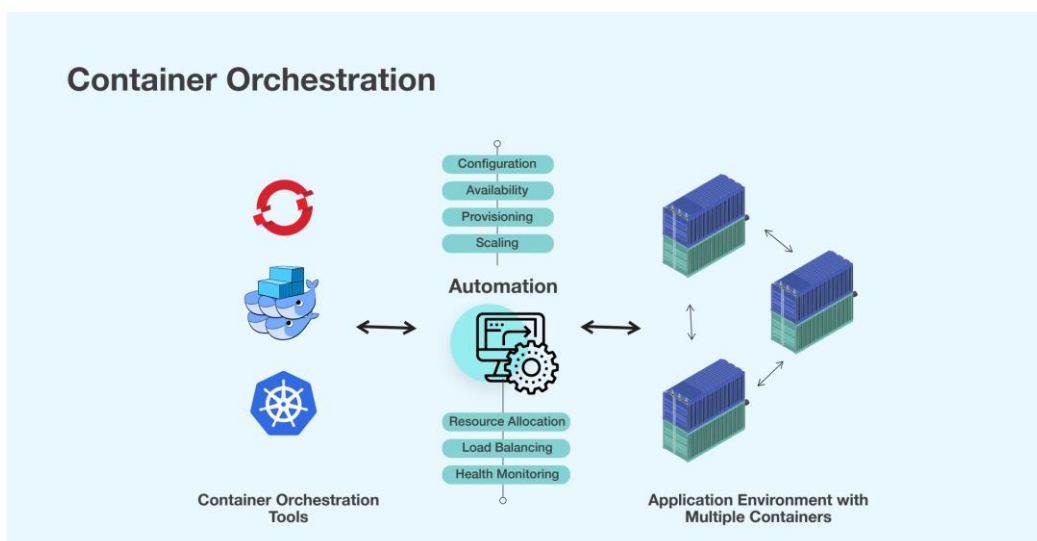
```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: userservice
5 spec:
6   replicas: 3
7   selector:
8     matchLabels:
9       app: userservice
10  template:
11    metadata:
12      labels:
13        app: userservice
14    spec:
15      containers:
16        - name: contactservice
17          image: 863099/springboot-image:latest
18          ports:
19            - containerPort: 3000
20
```

In a complete microservices architecture, we would have separate Deployment YAML files for each microservice, allowing us to independently manage and scale each service based on its specific requirements.

DEPLOYMENT STEPS:

- Containerizing Microservices with Docker
- Building Docker Images
- Pushing Docker Images to Container Registry
- Deploying Microservices with Kubernetes
- Creating Kubernetes Deployments using commands:
 - **kubectl apply -f deployment-service1.yaml**
 - **kubectl apply -f deployment-service2.yaml**
- Creating Kubernetes Services using commands:
 - **kubectl apply -f service-service1.yaml**
 - **kubectl apply -f service-service2.yaml**

Describe how container orchestration simplifies deployment, scaling, and management.



Container orchestration is the process of automating the operational effort required to run containerized workloads and services. It automates various aspects of the containers' lifecycle, including provisioning, deployment, scaling, networking, load balancing, traffic routing, and more.

Container orchestration simplifies deployment, scaling, and management by automating the processes of deploying and managing containers at scale. It abstracts complexities, streamlines resource allocation, automates scaling based on demand, and provides centralized management, making it easier to maintain, update, and scale applications in a containerized environment.

Deployment Process

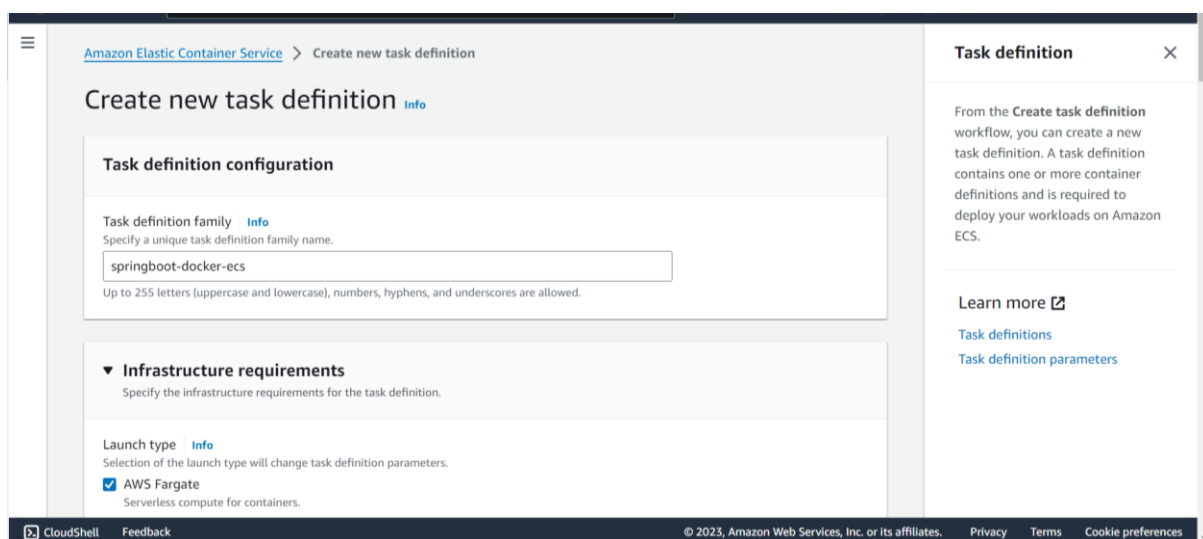
[Provide scripts or configurations for deploying the microservices on the chosen cloud platform.](#)

Prerequisites:

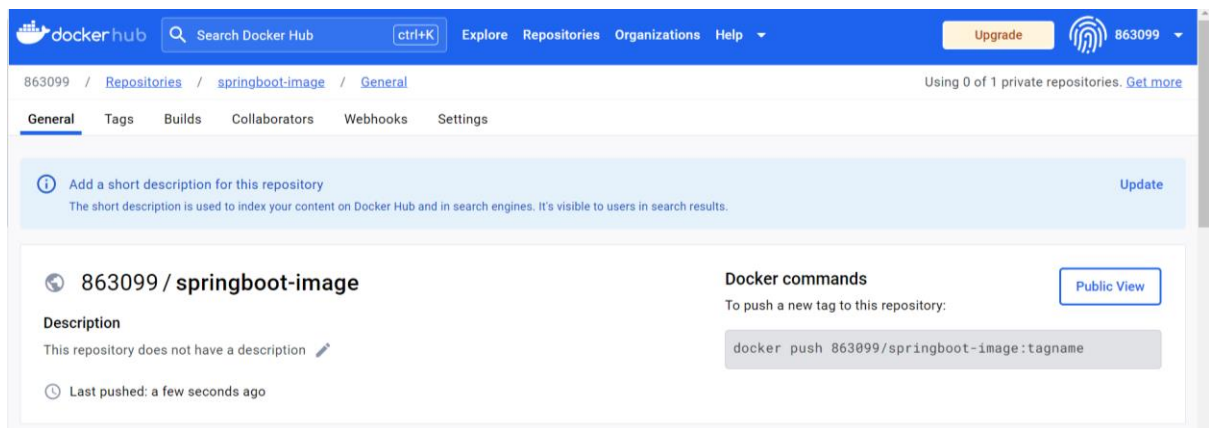
- Docker for Containerization (Already done)
- Kubernetes for Orchestration (Already done)
- Cloud Platform- AWS (Needs to be done)

WORKING ON AWS ECS AND AWS FARGATE:

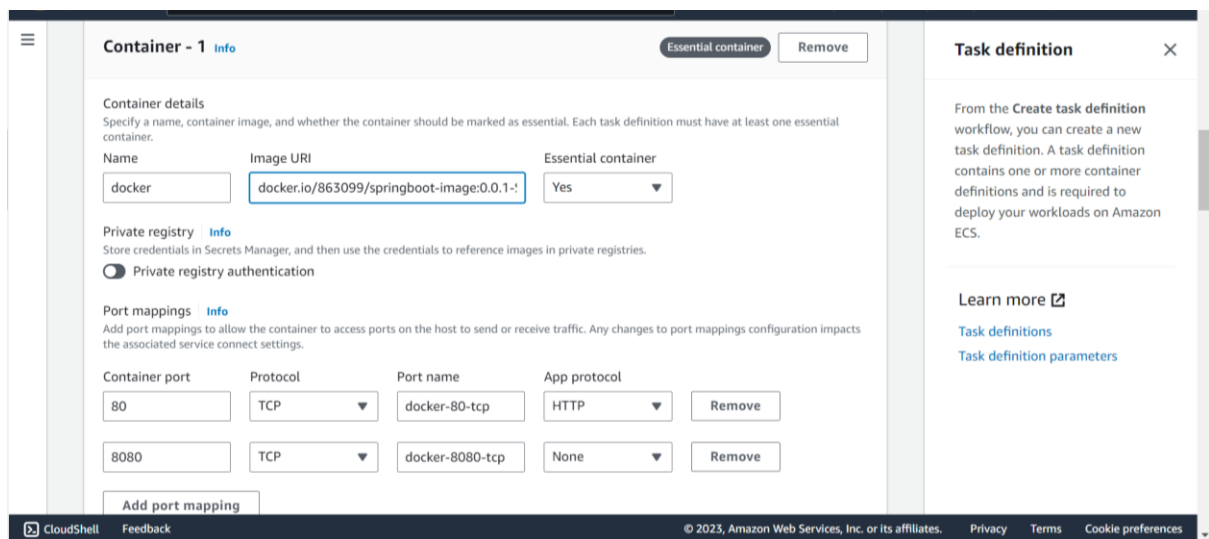
- In AWS ECS, creating task definition (Task definition contains one or more container definitions and is required to deploy our workload on AWS ECS)

The screenshot shows the 'Create new task definition' page in the AWS Management Console. The breadcrumb navigation at the top reads 'Amazon Elastic Container Service > Create new task definition'. The main heading is 'Create new task definition' with an 'Info' link. The page is divided into two main sections: 'Task definition configuration' and 'Infrastructure requirements'. In the 'Task definition configuration' section, there is a 'Task definition family' field with the value 'springboot-docker-ecs' entered. Below this field is a note: 'Up to 255 letters (uppercase and lowercase), numbers, hyphens, and underscores are allowed.' The 'Infrastructure requirements' section is expanded, showing 'Launch type' with the 'AWS Fargate' option selected. A note below the launch type says 'Selection of the launch type will change task definition parameters.' On the right side of the console, there is a 'Task definition' sidebar with a close button (X). It contains a description: 'From the Create task definition workflow, you can create a new task definition. A task definition contains one or more container definitions and is required to deploy your workloads on Amazon ECS.' Below this, there are links for 'Learn more', 'Task definitions', and 'Task definition parameters'. The footer of the console shows 'CloudShell', 'Feedback', '© 2023, Amazon Web Services, Inc. or its affiliates.', 'Privacy', 'Terms', and 'Cookie preferences'.

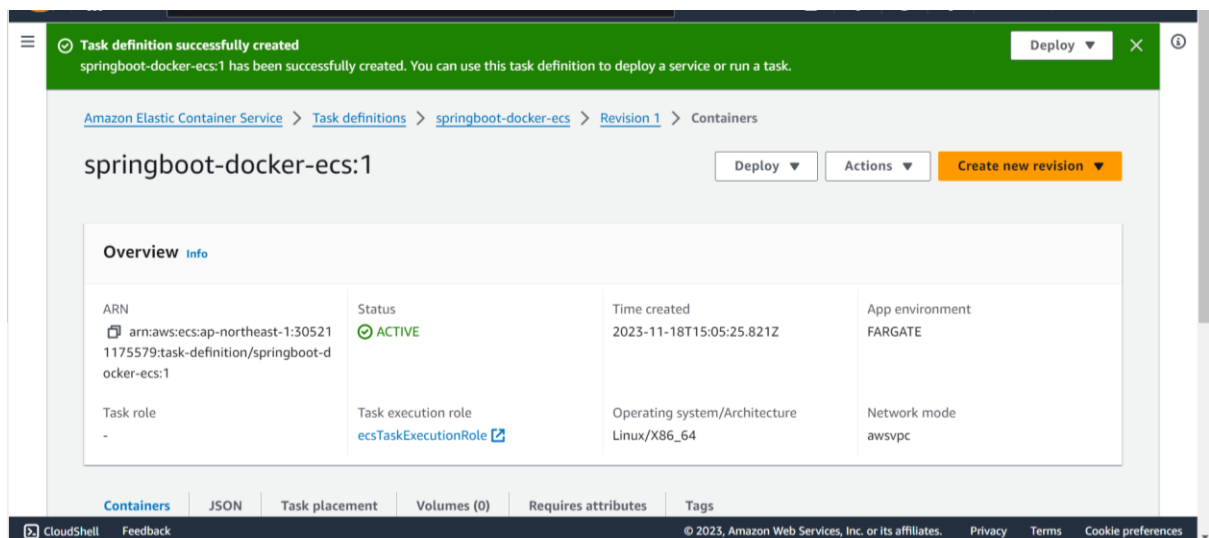
- Docker repository of image created on Docker Hub

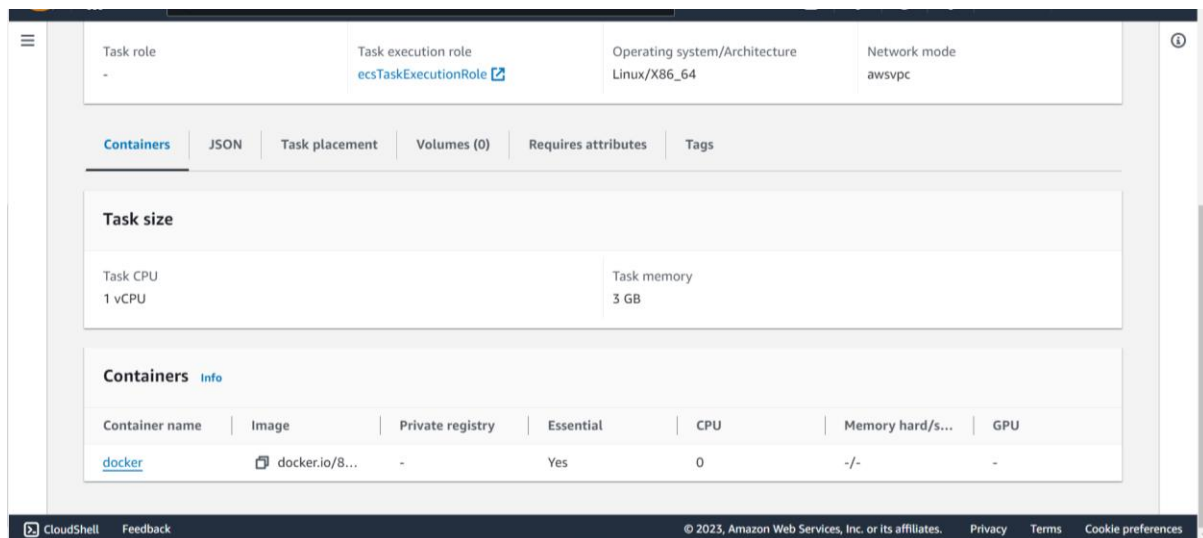


- Adding container and image url from docker hub

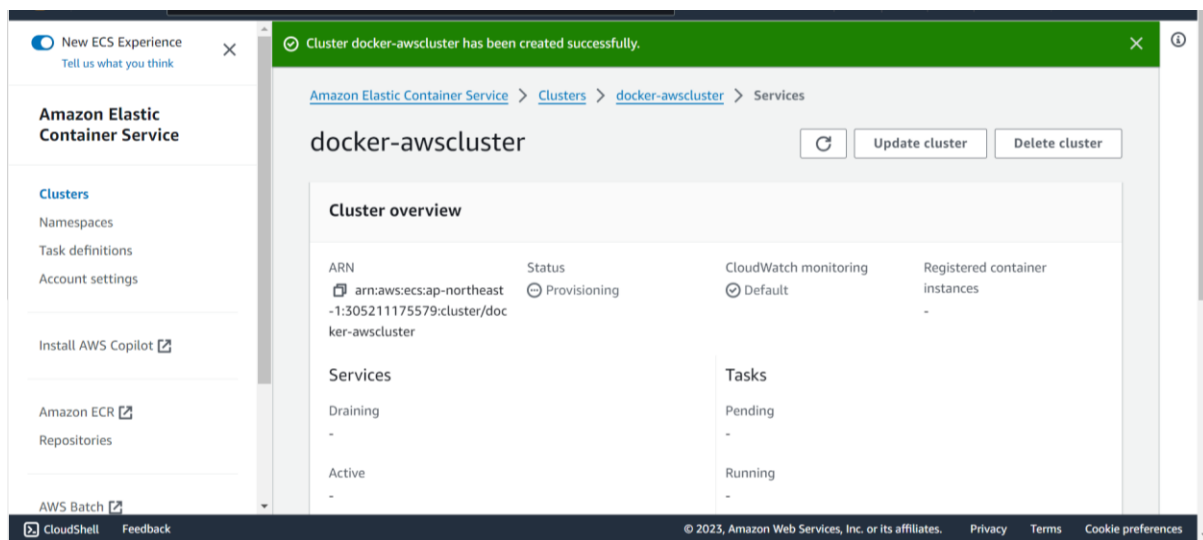


- Task definition created successfully

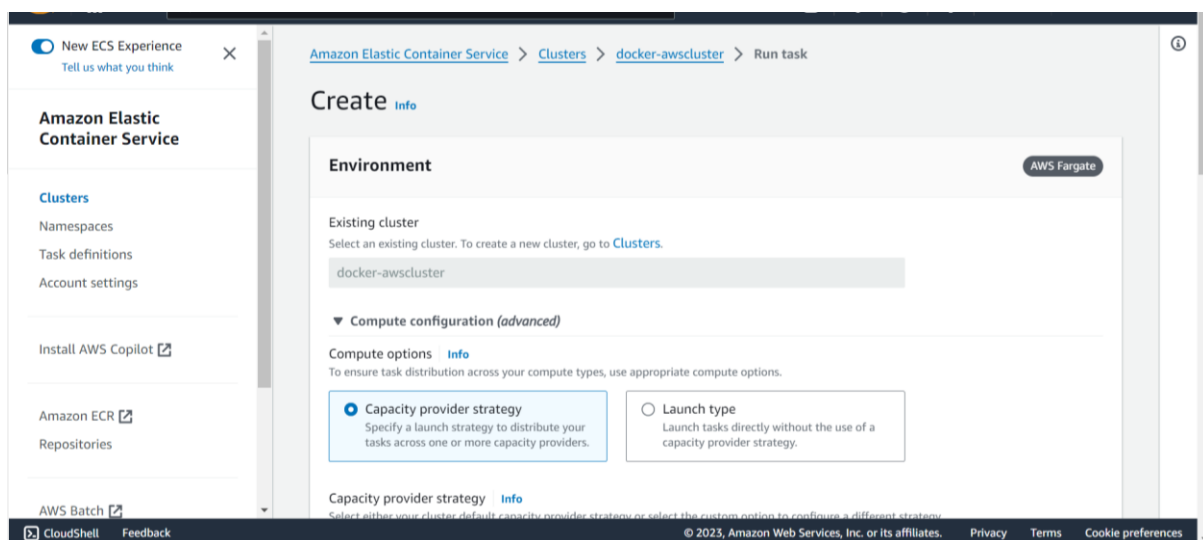




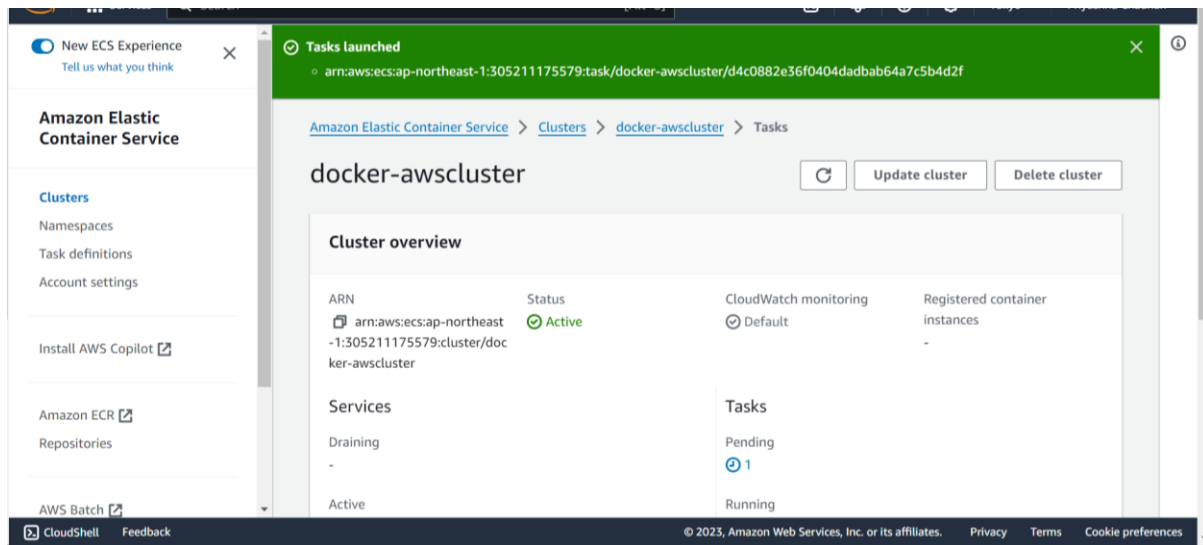
- Successfully Created cluster in AWS ECS.



- Adding Existing cluster to the task definition



- Task Launched successfully




Demonstrate the ability to deploy the application in a cloud environment consistently.

AWS provides a range of services that support microservices architecture:

AWS ECS

- Amazon Elastic Container Service.
- ECS is a fully managed container orchestration service.
- Amazon ECS makes it easy to deploy, manage and scale Docker containers running applications, services and batch processes.
- It places containers across our cluster based on our resource needs and is integrated with familiar features like Elastic Load Balancing, EC2 security groups, EBS volumes and IAM roles.



AWS Fargate

- AWS Fargate is a serverless compute engine for containers that works with ECS & it allow you to run containers without having provision, configure & Scaling.
- Eliminate the need of EC2 instance.
- AWS Fargate is compatible with both Amazon Elastic Container Service (Amazon ECS) and Amazon Elastic Kubernetes Service (Amazon EKS).

ENTIRE PROCESS:

Step 1: Developing Spring Boot Application- Creating Spring Boot application

Step 2: Dockerizing the Spring Boot Application- Writing a Dockerfile and building Docker images locally

Step 3: Building and Testing Docker Image Locally

Step 4: Pushing Docker Image to Docker Hub


Step 5: Setting up AWS ECS CLUSTER- Creating an ECS Cluster on AWS

Step 6: Creating ECS Task Definition- Defining task definition specifying the Docker image.

Step 7: Creating ECS service- Setting up an ECS service using task definition

Step 8: Configuring AWS Fargate- Choosing Fargate as the launch type.

Step 9: Deploying on AWS Fargate- Deploying the ECS service and Fargate will manage the containers.



Scaling and Load Balancing

Implement automatic scaling mechanisms for the microservices based on resource usage.

Kubernetes provides Horizontal Pod Autoscaling (HPA) for automatically adjusting the number of replicas in a deployment based on observed CPU or memory utilization.

Add Resource Requests and Limits to your Deployment:

Updating Kubernetes deployment YAML files to include resource requests and limits for CPU and memory:

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: contactservice
5 spec:
6   replicas: 3
7   selector:
8     matchLabels:
9       app: contactservice
10  template:
11    metadata:
12      labels:
13        app: contactservice
14    spec:
15      containers:
16      - name: contactservice
17        image: 863099/springboot-image:latest
18        ports:
19        - containerPort: 3000
20      resources:
21        requests:
22          memory: "64Mi"
23          cpu: "250m"
24        limits:
25          memory: "128Mi"
26          cpu: "500m"
```

Add Horizontal Pod Autoscaler (HPA):

Creating an HPA for each microservice to automatically adjust the number of replicas based on CPU or memory usage.

```
1 apiVersion: autoscaling/v2beta2
2 kind: HorizontalPodAutoscaler
3 metadata:
4   name: contactservice-hpa
5 spec:
6   scaleTargetRef:
7     apiVersion: apps/v1
8     kind: Deployment
9     name: contactservice
10  minReplicas: 3
11  maxReplicas: 10
12  metrics:
13  - type: Resource
14    resource:
15      name: cpu
16      target:
17        type: Utilization
18        averageUtilization: 50
```

Applying the updated deployment files and the new HPA files to Kubernetes cluster using commands:


kubectl apply -f deployment-service1.yaml

kubectl apply -f hpa-service1.yaml

Monitoring the scaling activities using the following command:

kubectl get hpa

[Explain how load balancing is achieved within the containerized environment.](#)



Load balancing is a technique used to distribute incoming network traffic across multiple servers or instances to enhance availability, scalability, and performance. In containerized environments like Kubernetes, load balancing ensures even distribution of requests among containers, optimizing resource utilization and allowing for uninterrupted performance of microservices-based applications. It enables updating or scaling individual microservices without disruptions by distributing traffic effectively across the containerized services. Load balancing is critical for achieving high availability and reliability in distributed systems.

Monitoring and Logging

[Set up monitoring and logging for the deployed microservices.](#)

For Setting up monitoring and logging for deployed microservices, used popular tools such as Prometheus for monitoring and Elasticsearch-Fluentd-Kibana (EFK) stack for logging in a Kubernetes environment.

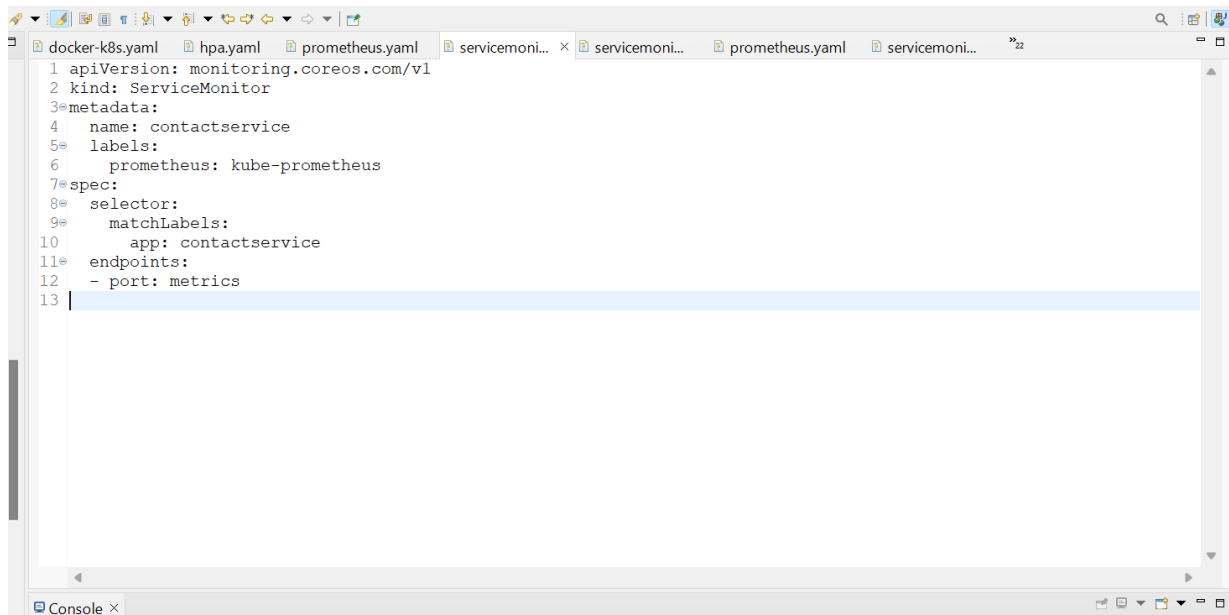
Monitoring with Prometheus:

```
docker-k8s.yaml | hpa.yaml | prometheus.yaml x | prometheus.yaml | servicemoni... | prometheus.yaml | servicemoni...
prometheus.json (prometheus.json)
1 apiVersion: monitoring.coreos.com/v1
2 kind: Prometheus
3 metadata:
4   name: prometheus
5 spec:
6   replicas: 1
7   ruleSelector:
8     matchLabels:
9       prometheus: kube-prometheus
10  serviceMonitorSelector:
11    matchLabels:
12      prometheus: kube-prometheus
13  resources:
14    requests:
15      memory: 400Mi
16  alerting:
17    alertmanagers:
18      - namespace: default
19        name: alertmanager-main
20        port: web
```

Applying the configuration using command:

kubectl apply -f prometheus.yaml

Setting Up Service Monitors:

A screenshot of a code editor window with multiple tabs. The active tab is 'prometheus.yaml'. The code is a YAML configuration for a ServiceMonitor. The tabs include 'docker-k8s.yaml', 'hpa.yaml', 'prometheus.yaml', 'servicemoni...', 'servicemoni...', 'prometheus.yaml', 'servicemoni...', and '»22'. The code is as follows:

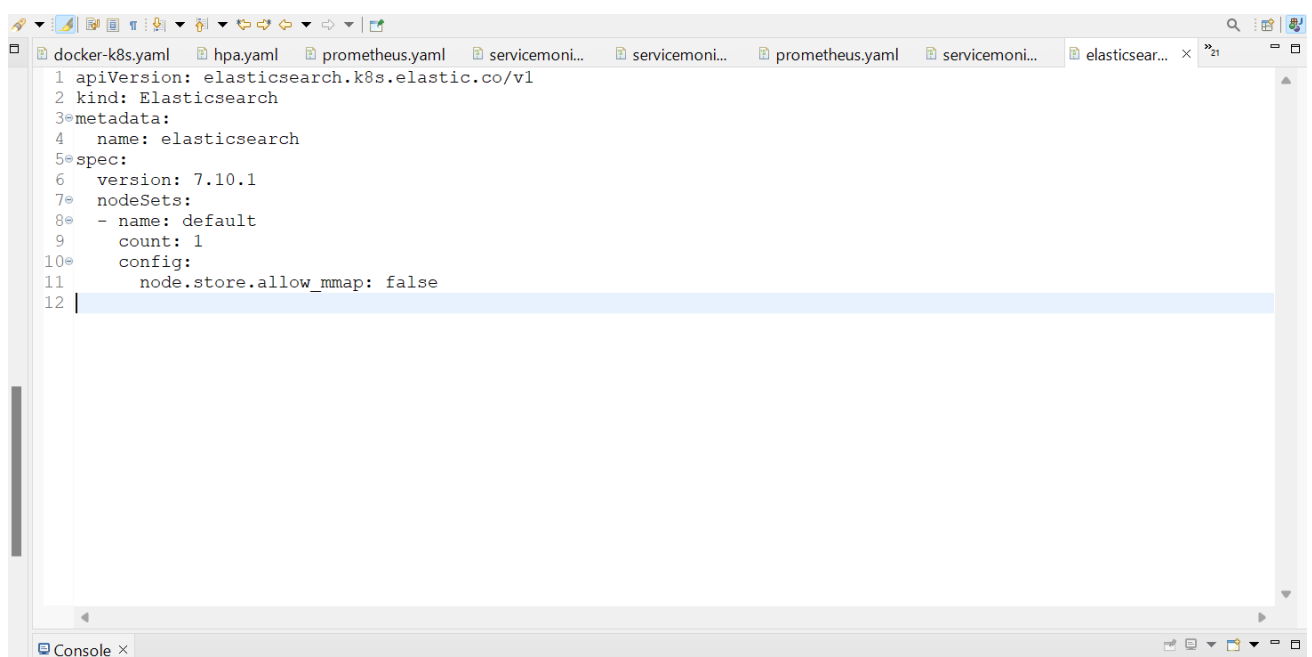
```
1 apiVersion: monitoring.coreos.com/v1
2 kind: ServiceMonitor
3 metadata:
4   name: contactservice
5   labels:
6     prometheus: kube-prometheus
7 spec:
8   selector:
9     matchLabels:
10      app: contactservice
11 endpoints:
12   - port: metrics
13
```

Applying the configuration using command:

kubectl apply -f servicemonitor-service1.yaml

Logging with EFK (Elasticsearch, Fluentd, Kibana):

Deploying Elasticsearch:

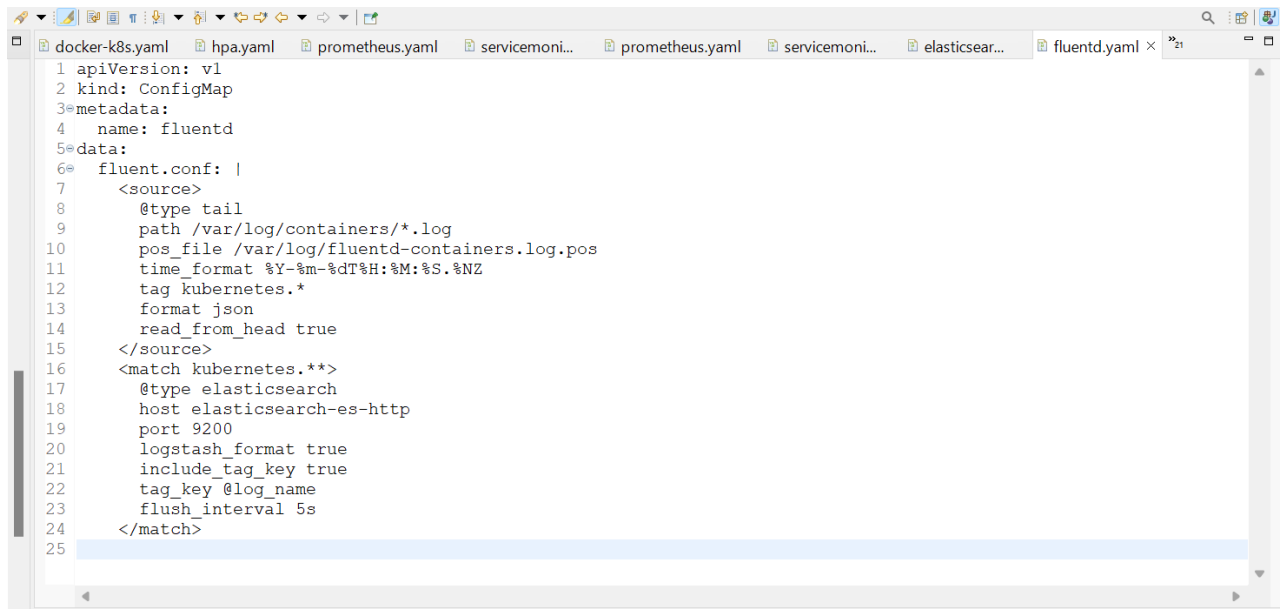
A screenshot of a code editor window with multiple tabs. The active tab is 'elasticsearch...'. The code is a YAML configuration for an Elasticsearch deployment. The tabs include 'docker-k8s.yaml', 'hpa.yaml', 'prometheus.yaml', 'servicemoni...', 'servicemoni...', 'prometheus.yaml', 'servicemoni...', and 'elasticsearch...'. The code is as follows:

```
1 apiVersion: elasticsearch.k8s.elastic.co/v1
2 kind: Elasticsearch
3 metadata:
4   name: elasticsearch
5 spec:
6   version: 7.10.1
7   nodeSets:
8     - name: default
9       count: 1
10      config:
11        node.store.allow_mmap: false
12
```

Applying the configuration using command:

kubectl apply -f elasticsearch.yaml

Deploying Fluentd:

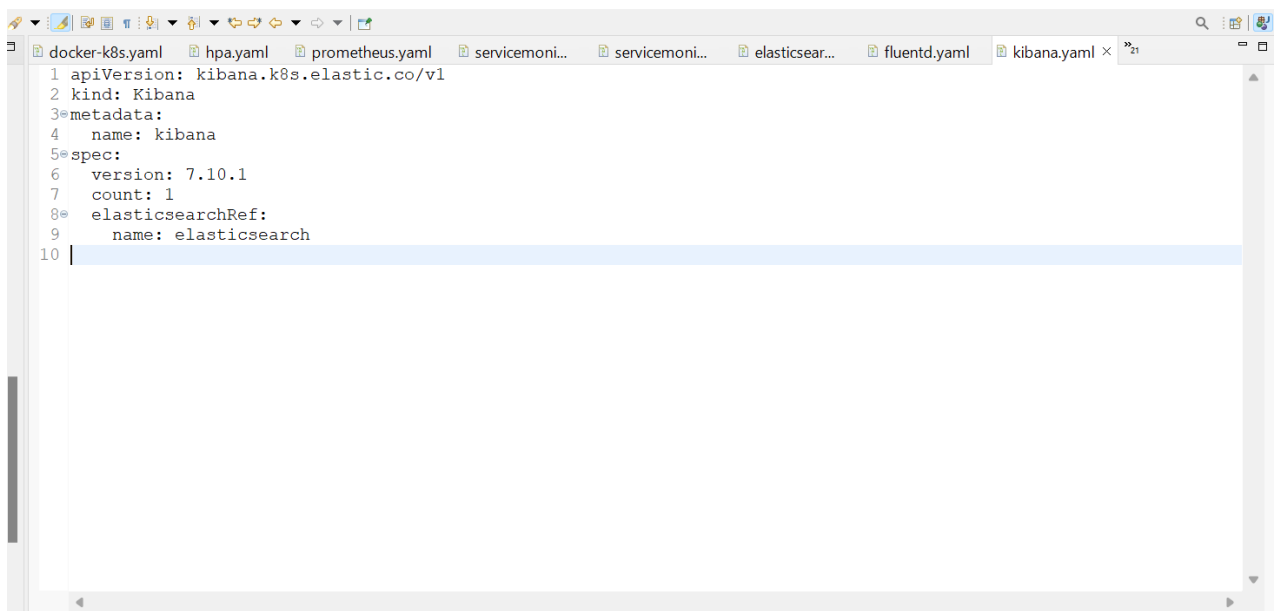
A screenshot of a code editor window with multiple tabs. The active tab is 'fluentd.yaml'. The code is a Kubernetes ConfigMap definition for Fluentd. It specifies the source as tailing logs from /var/log/containers and the sink as Elasticsearch at port 9200. The configuration is enclosed in XML-like tags for source and match.

```
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: fluentd
5 data:
6   fluent.conf: |
7     <source>
8       @type tail
9       path /var/log/containers/*.log
10      pos_file /var/log/fluentd-containers.log.pos
11      time_format %Y-%m-%dT%H:%M:%S.%NZ
12      tag kubernetes.*
13      format json
14      read_from_head true
15    </source>
16    <match kubernetes.**>
17      @type elasticsearch
18      host elasticsearch-es-http
19      port 9200
20      logstash_format true
21      include_tag_key true
22      tag_key @log_name
23      flush_interval 5s
24    </match>
25
```

Applying the configuration using command:

kubectl apply -f fluentd.yaml

Deploying Kibana:

A screenshot of a code editor window with multiple tabs. The active tab is 'kibana.yaml'. The code is a Kubernetes Deployment definition for Kibana. It specifies the image as kibana.k8s.elastic.co/v1, version 7.10.1, and references the elasticsearch service.

```
1 apiVersion: kibana.k8s.elastic.co/v1
2 kind: Kibana
3 metadata:
4   name: kibana
5 spec:
6   version: 7.10.1
7   count: 1
8   elasticsearchRef:
9     name: elasticsearch
10
```

Applying the configuration using command:

kubectl apply -f kibana.yaml

Showcase the ability to monitor the application's performance and troubleshoot issues.

Monitoring:

- Deploying Prometheus for metrics collection.
- Visualizing metrics using Grafana dashboards.
- Setting up alerts for critical metrics.

Troubleshooting:

- Using `kubectl logs` to inspecting container logs.
- Monitoring resource utilization with `kubectl top`.
- Implementing health probes for application reliability.
- Leveraging Kubernetes auto-scaling for dynamic scaling.
- Using distributed tracing for end-to-end visibility.
- Regularly audit access controls for security.
- Recreating pods to troubleshoot transient issues.